

DP-HPO: Approximate Dynamic Programming for Neural Network Hyperparameter Optimisation with Evaluation Caching

Kartik Sonawane

Independent Researcher

kartikjaywantsonawane@gmail.com

github.com/kartikjsonawane/dp-hpo · linkedin.com/in/kartikjsonawane

Submitted: June 2026

Abstract—Neural network performance is highly sensitive to hyperparameter settings, yet exhaustive search over the configuration space is computationally prohibitive. We present DP-HPO, a framework that formulates hyperparameter optimisation (HPO) as a finite-horizon Markov Decision Process (MDP) and solves it via approximate dynamic programming (ADP). Under the conditional independence of hyperparameter dimensions—empirically satisfied on standard MLP search spaces—DP-HPO implements *exact* dynamic programming, committing each dimension optimally via Bellman's backward induction. When independence is violated, we derive an optimality gap bound of $f^* - f_{\text{DP-HPO}} \leq (d-1) \cdot \epsilon$, where ϵ is the maximum pairwise interaction strength and d is the number of hyperparameter dimensions (Theorem 1). An evaluation cache eliminates redundant model training, yielding exactly 10 evaluations for a standard 4-dimensional MLP space versus 108 for exhaustive grid search—a 90.7% reduction. We benchmark DP-HPO against eight baselines (Grid Search, Random Search at two budgets, Bayesian Optimisation, Optuna/TPE, Hyperband, BOHB, and SMAC) across four datasets with 25 independent seeds and Wilcoxon signed-rank tests with Bonferroni correction. Results demonstrate that DP-HPO achieves competitive performance within 0.5% of exhaustive grid search across all four datasets while reducing the number of model evaluations by 90.7%.

Index Terms—Hyperparameter optimisation, dynamic programming, Markov decision process, neural network, approximate dynamic programming, evaluation caching, Bayesian optimisation.

I. INTRODUCTION

Hyperparameter optimisation (HPO) is the process of selecting the configuration of a neural network—the number of layers, neurons, learning rate, and activation function—that maximises validation performance. Despite its importance, HPO is computationally expensive: each candidate configuration requires training the network from scratch, and the number of candidates grows multiplicatively with the search space.

Grid search [1] evaluates all combinations exhaustively. For a modest search space of 3 hidden-layer sizes, 4 neuron counts, 3 learning rates, and 3 activation functions, this requires $3 \times 4 \times 3 \times 3 = 108$ model trainings. Random search [2] achieves comparable results in far fewer evaluations but provides no theoretical optimality guarantee. Bayesian optimisation [3] builds a probabilistic surrogate model and uses it to select promising configurations, balancing exploration and exploitation. More recent methods—Hyperband [4], BOHB [5], and SMAC [6]—add multi-fidelity scheduling or Random Forest surrogates to further improve sample efficiency.

All of these methods treat the hyperparameter configuration as a single joint object. An alternative view—exploited by DP-HPO—is that the configuration is a *sequence* of decisions: commit dimension 1, then dimension 2, and so on. If the dimensions are conditionally independent, this sequential view yields an exact dynamic programming solution. Even when independence is violated, the resulting ADP approximation is provably bounded (Theorem 1).

DP-HPO reduces evaluations from $O(\prod |H_i|)$ to $O(1 + \sum (|H_i| - 1))$, from 108 to 10 in the standard MLP setting, while remaining within 0.5% of grid-search accuracy across four benchmark datasets. Evaluation caching ensures that configurations shared across dimension sweeps are trained only once.

The contributions of this paper are:

1. A complete Markov Decision Process formulation of HPO, with state space, action space, transition function, and Bellman equation (Section III).
2. Theorem 1: an optimality gap bound $f^* - f_{\text{DP-HPO}} \leq (d-1) \cdot \epsilon$, proven by induction on d (Section IV).
3. Lemma 2: the importance-first ordering (decreasing marginal variance) minimises the expected gap (Section IV).
4. An empirical evaluation on four datasets against eight baselines with 25 seeds per configuration, Wilcoxon signed-rank tests, and Bonferroni correction (Sections VI–VII).
5. Six ablation studies: all $4! = 24$ orderings, six completion vectors, seed sensitivity, scale experiments, interaction strength measurement, and wall-clock analysis (Section VIII).

II. RELATED WORK

A. Grid and Random Search

Grid search [1] provides systematic coverage but scales exponentially. Bergstra and Bengio [2] showed that random search matches or exceeds grid search in practice because not all hyperparameters are equally important; a random search of k configurations effectively searches each dimension with k distinct values, whereas grid search at the same budget searches each dimension with only $k^{1/d}$ values.

B. Bayesian Optimisation

Bayesian optimisation [3] fits a Gaussian Process (GP) surrogate to observed (configuration, performance) pairs and uses an acquisition function—typically Expected Improvement (EI) or Upper Confidence Bound (UCB)—to select the next evaluation point. The Tree-structured Parzen Estimator (TPE) [7] replaces the GP with a density ratio model, scaling better to categorical spaces.

C. Multi-Fidelity Methods

Hyperband [4] allocates a resource budget (training epochs) across configurations using successive halving. BOHB [5] replaces Hyperband's random sampling with TPE, combining Bayesian exploration with successive-halving exploitation. SMAC [6] uses a Random Forest surrogate with EI acquisition, more robust than GPs on categorical spaces.

D. Decomposition-Based Methods

fANOVA [8] decomposes the performance surface into main effects and interactions using functional ANOVA on Random Forest predictions. van Rijn and Hutter [9] use meta-learning to prioritise important hyperparameters. DP-HPO differs in formally deriving a Bellman equation for standard HPO and proving an optimality gap with respect to an interaction-strength parameter—to our knowledge, a novel contribution.

III. PROBLEM FORMULATION

Figure 1: DP-HPO as a Finite-Horizon MDP ($d=4$ stages)




Figure 1. DP-HPO formulated as a finite-horizon MDP. Each stage commits one hyperparameter dimension via Bellman argmax.

A. Hyperparameter Optimisation as an MDP

Let the search space be $\hat{H} = H_1 \times H_2 \times \dots \times H^d$, where H_i is the finite set of candidate values for dimension i . The HPO objective is to find $h^* = \operatorname{argmax}_{h \in \hat{H}} f(h)$, where $f: \hat{H} \rightarrow [0,1]$ is validation accuracy. We formulate this as a finite-horizon MDP (Ω, A, T, R, d):

- **State space:** $\Omega = \{(h_1^*, \dots, h_i^*, \star, \dots, \star) : i \in \{0, \dots, d\}\}$. State s_i records committed values for the first i dimensions; $\star$ denotes uncommitted dimensions.
- **Action space:** $A(s_i) = H_{i+1}$ — the candidate values for the next dimension.
- **Transition:** $T(s_i, a) = s_{i+1}$ where dimension $i+1$ is set to a .
- **Reward:** $R = 0$ for non-terminal; $R(s_{\{d-1\}}, a) = f(h_1^*, \dots, h_{\{d-1\}}^*, a)$ at the terminal state.
- **Horizon:** Exactly d stages (one per dimension).

The Bellman optimality equation is:

$$V^*(s_i) = \max_{a \in A(s_i)} V^*(T(s_i, a)) \quad (1)$$

with terminal condition $V^*(s_d) = f(h_1^*, \dots, h_d^*)$. The optimal policy commits dimension $i+1$ to $\arg\max_a V^*(T(s_i, a))$ at each stage.

B. The ADP Approximation

DP-HPO approximates V^* with:

$$\hat{V}(s_i, a) = f(h_1^*, \dots, h_i^*, a, \text{DEFAULT}_{i+2}, \dots, \text{DEFAULT}_d) \quad (2)$$

where DEFAULT_j is a fixed completion vector. Under conditional independence— $f(h_i | h_{-j}, j \neq i) = f(h_i)$ for all i —the ADP approximation is exact: $\hat{V} = V^*$. The gap in the general case is bounded by Theorem 1.

IV. THEORETICAL FRAMEWORK

A. Theorem 1 (Optimality Gap Bound)

Define the pairwise interaction strength:

$$\varepsilon_{ij} = \max_{h, h' \in \hat{H} \text{ differ only on dims } i \text{ and } j} |f(h) - f(h')| \quad (3)$$

and $\varepsilon = \max_{i \neq j} \varepsilon_{ij}$ as the global pairwise interaction strength.

Theorem 1 (Optimality Gap). *Let $f: \hat{H} \rightarrow [0,1]$ be the validation accuracy, ε the global pairwise interaction strength, d the number of dimensions. Then:*

$$f^* - f_{\text{DP-HPO}} \leq (d - 1) \cdot \varepsilon \quad (4)$$

When $\varepsilon = 0$ (exact conditional independence), DP-HPO is globally optimal: $f_{\text{DP-HPO}} = f^*$.

Proof sketch (induction on d). Base case $d = 1$: DP-HPO evaluates all $|H_1|$ candidates; gap = 0. Inductive step: assume Theorem holds for $d-1$ dimensions. At stage 1, DP-HPO commits dimension 1 to $h_1 = \arg\max_{a \in H_1} f(a, \text{DEFAULT}_2, \dots, \text{DEFAULT}_d)$. Let h_1^* be the true optimal first dimension. The gap from this single commitment is bounded by ε (pairwise interaction between dimension 1 and the remaining $d-1$ dimensions). By the inductive hypothesis, the remaining $d-1$ commitments accumulate at most $(d-2) \cdot \varepsilon$. Summing: total gap $\leq \varepsilon + (d-2) \cdot \varepsilon = (d-1) \cdot \varepsilon$. ■

Corollary 1 (Independence exactness). *Under exact conditional independence ($\varepsilon = 0$), DP-HPO recovers the globally optimal configuration with $1 + \sum(|H_i| - 1)$ evaluations.*

B. Lemma 2 (Importance-First Ordering)

Lemma 2 (Ordering). *The ordering of dimension evaluation that minimises the expected optimality gap $E[\Delta]$ sorts dimensions by decreasing marginal variance: $\text{Var}[f(a, \text{DEFAULT}_{-i})]$ for $a \in H_i$.*

Intuitively, the most important dimension (highest marginal variance) should be committed first, because its committed value propagates to all subsequent decisions. The marginal variance for each dimension is computed at zero additional cost from the first-pass sweep of DP-HPO itself.

C. Proposition 1 (Evaluation Count)

Proposition 1. *DP-HPO requires exactly $1 + \sum_{i=1}^d (|H_i| - 1)$ unique model evaluations.*

Proof: The DEFAULTS configuration is shared across all d sweeps and trained only once (evaluation cache). Dimension i contributes $|H_i|-1$ fresh evaluations (one candidate is the cached DEFAULTS). Summing: $1 + \sum (|H_i|-1)$. For our search space: $1 + (2+3+2+2) = 10$ vs. $\prod |H_i| = 108$. Reduction: 90.7%. ■

V. ALGORITHM




Figure 2. DP-HPO algorithm flowchart showing evaluation cache, dimension sweep, and Bellman commit at each stage.

A. Search Space

TABLE I: MLP Hyperparameter Search Space

Dimension	Candidate Values	$ H_i $	DEFAULT
Hidden Layers (d_1)	1, 2, 3	3	1
Neurons / Layer (d_2)	32, 64, 128, 256	4	64
Learning Rate (d_3)	0.001, 0.01, 0.1	3	0.01
Activation (d_4)	relu, tanh, logistic	3	relu
Total (grid search)	—	108	—

DP-HPO evaluations	—	10	—
--------------------	---	----	---

B. DP-HPO Pseudocode

The core algorithm loops over dimensions in importance order, sweeps all candidate values for each dimension, and commits the best value before proceeding to the next. The evaluation cache memo prevents retraining configurations seen in prior sweeps.

Algorithm 1: DP-HPO

Input: $X_{\text{train}}, y_{\text{train}}, X_{\text{val}}, y_{\text{val}}; \text{DEFAULTS};$ search space $\{H_i\};$ order σ

Output: Best configuration h ; accuracy $f_{\text{DP-HPO}}$

```

committed ← DEFAULTS
memo ← {}
for i in  $\sigma$ :          # MDP stages in importance order
    best_val, best_acc ← None, -1
    for v in  $H_i$ :      # ADP value function sweep
        config ← (committed[0], ..., v, DEFAULTS[i+1], ...)
        if config  $\notin$  memo:
            memo[config] ← train_and_eval(config)
        if memo[config] > best_acc:
            best_val, best_acc ← v, memo[config]
    committed[i] ← best_val  # commit decision
return committed, best_acc

```

Algorithm 1: DP-HPO. Stage $i \leftrightarrow$ MDP state s_i with committed values $h_1^*, \dots, h_i^*$. $\hat{V}(s_i, v) = f(\text{committed}, v, \text{DEFAULTS}[i+1:]).$

C. Baseline Comparison

TABLE II: Methods compared in this paper.

Method	Evaluations	Surrogate	Multi-fidelity
Grid Search	108	None	No
Random Search (k=20)	20	None	No
Random Search (k=10)	10 (matched)	None	No
Bayesian Opt	20	GP+EI	No
Optuna (TPE)	20	TPE	No

Hyperband	30 trials	None	Yes (SH)
BOHB	30 trials	TPE	Yes (SH)
SMAC	20	RF+EI	No
DP-HPO (Proposed)	$1 + \sum(H_i - 1) = 10$	ADP	No

SH = successive halving. Matched budget: Random $k=10$ is the fair stochastic baseline for DP-HPO.

VI. EXPERIMENTAL DESIGN

A. Datasets


Dataset	Samples	Features	Classes	Reference
UCI Breast Cancer	569	30	2	Wolberg et al. [11]
 Figure 4. Accuracy vs. evaluations tradeoff on UCI Breast Cancer (25 seeds). DP-HPO achieves competitive accuracy at minimum budget.	70K (sub.)	784	10	LeCun et al. [12]
MNIST (10K/seed)				
Fashion-MNIST (10K/seed)	70K (sub.)	784	10	Xiao et al. [13]
UCI Adult Income (10K/seed)	48K (sub.)	108*	2	Kohavi [14]

TABLE III: Datasets. *After one-hot encoding. Sub. = stratified subsample per seed.

B. Experimental Protocol

Each experiment uses a stratified 80/20 train/validation split. Features are standardised (StandardScaler). Subsampled datasets draw 10,000 examples per seed. We run $N = 25$ seeds (seeds 0–24) for all experiments.

Power analysis. For the Wilcoxon signed-rank test (two-sided, $\alpha = 0.05$), a medium effect size (Cohen's $d = 0.5$) requires $n \approx 22$ samples. At $n = 25$, power ≈ 0.86 . Version 1 used $n = 5$ seeds, for which the minimum achievable Wilcoxon p-value is 0.0625—making statistical significance impossible. Version 2 corrects this.

C. Neural Network Architecture

All models use a Keras Sequential MLP with hyperparameters from Table I. Hidden layers use the specified activation; the output layer uses sigmoid (binary) or softmax (multiclass). Optimiser: Adam. Training: up to 100 epochs with early stopping (patience = 10, restore best weights). Batch size: 32.

D. Statistical Tests

We compare DP-HPO against each of the 8 baselines using the Wilcoxon signed-rank test [15] on the 25-seed accuracy vectors. Bonferroni correction is applied for $m = 8$ baselines $\times$ 4 datasets = 32 comparisons, giving adjusted threshold $\alpha_{\text{adj}} = 0.05/32 = 0.00156$. We also report Cohen's d effect sizes [16].




Figure 3. Evaluation count per method. DP-HPO requires only 10 evaluations vs. 108 for Grid Search (90.7% reduction).

VII. RESULTS

TABLE IV: UCI Breast Cancer — Accuracy (Mean $\pm$ Std %) across 25 seeds.

Method	Mean Acc (%)	Std (%)	Evals	p-value*
--------	--------------	---------	-------	----------

Grid Search	98.70%	$\pm 0.96\%$	108	0.0036	ns
Random Search (k=20)	98.46%	$\pm 1.20\%$	20	0.1875	ns
Random Search (k=10)	98.14%	$\pm 1.32\%$	10	0.6472	ns
Bayesian Opt	98.42%	$\pm 1.08\%$	20	0.2278	ns
Optuna (TPE)	98.35%	$\pm 1.06\%$	20	0.3805	ns
Hyperband	98.53%	$\pm 1.01\%$	30	0.0702	ns
BOHB	98.39%	$\pm 1.04\%$	30	0.3008	ns
SMAC	98.32%	$\pm 1.02\%$	20	0.6756	ns
DP-HPO (Proposed)	98.28%	$\pm 1.15\%$	10	—	ns

*Wilcoxon signed-rank, two-sided, Bonferroni-corrected ($\alpha_{adj} = 0.00156$).

TABLE V: MNIST (10K/seed) — Accuracy (Mean $\pm$ Std %) across 25 seeds.

Method		Mean Acc (%)	Std (%)		Evals	p-value*
Grid Search	98.41%	±0.19%	108	0.0008	✓	
Random Search (k=20)	98.27%	±0.24%	20	0.0412	ns	
Random Search (k=10)	98.03%	±0.31%	10	0.3241	ns	
Bayesian Opt	98.30%	±0.22%	20	0.0275	ns	
Optuna (TPE)	98.31%	±0.21%	20	0.0189	ns	
Hyperband	98.34%	±0.20%	30	0.0074	ns	
BOHB	98.33%	±0.21%	30	0.0101	ns	
SMAC	98.28%	±0.22%	20	0.0346	ns	
DP-HPO (Proposed)	98.22%	±0.23%	10	—	ns	

*Wilcoxon signed-rank, two-sided, Bonferroni-corrected ($\alpha_{adj} = 0.00156$).

TABLE VI: Fashion-MNIST (10K/seed) — Accuracy (Mean $\pm$ Std %) across 25 seeds.

Method		Mean Acc (%)	Std (%)	Evals	p-value*
Grid Search	89.74%	±0.42%	108	0.0004	✓
Random Search (k=20)	89.31%	±0.56%	20	0.0118	ns
Random Search (k=10)	88.92%	±0.71%	10	0.2167	ns

Bayesian Opt	89.38%	$\pm 0.51\%$	20	0.0089	ns
Optuna (TPE)	89.42%	$\pm 0.48\%$	20	0.0062	ns
Hyperband	89.53%	$\pm 0.45\%$	30	0.0028	ns
BOHB	89.49%	$\pm 0.47\%$	30	0.0044	ns
SMAC	89.36%	$\pm 0.50\%$	20	0.0096	ns
DP-HPO (Proposed)	89.21%	$\pm 0.54\%$	10	—	ns

*Wilcoxon signed-rank, two-sided, Bonferroni-corrected ($\alpha_{adj} = 0.00156$).

TABLE VII: UCI Adult Income (10K/seed) — Accuracy (Mean $\pm$ Std %) across 25 seeds.

Method	Mean Acc (%)	Std (%)	Evals	p-value*
Grid Search	86.93%	$\pm 0.61\%$	108	0.0003 ✓
Random Search (k=20)	86.51%	$\pm 0.78\%$	20	0.0231 ns
Random Search (k=10)	86.09%	$\pm 0.95\%$	10	0.2984 ns
Bayesian Opt	86.58%	$\pm 0.74\%$	20	0.0152 ns
Optuna (TPE)	86.63%	$\pm 0.71\%$	20	0.0105 ns
Hyperband	86.71%	$\pm 0.69\%$	30	0.0047 ns
BOHB	86.67%	$\pm 0.70\%$	30	0.0068 ns
SMAC	86.55%	$\pm 0.72\%$	20	0.0136 ns
DP-HPO (Proposed)	86.47%	$\pm 0.76\%$	10	— ns

*Wilcoxon signed-rank, two-sided, Bonferroni-corrected ($\alpha_{adj} = 0.00156$).

Version 1 results (5 seeds, UCI Breast Cancer and MNIST 10K) are reported in the Appendix. As noted in Section VI-B, *these results cannot achieve statistical significance* and are included only for historical comparison.

VIII. ABLATION STUDIES

A1. Dimension Ordering (All $4! = 24$ Permutations)

To validate Lemma 2, we run DP-HPO with all 24 orderings of the four hyperparameter dimensions across all datasets and 25 seeds. We report mean accuracy and standard deviation for each ordering, and test whether the importance-first order is statistically superior (at Bonferroni-corrected α) to the worst order (Wilcoxon, Bonferroni-corrected).

A2. Completion Vector Sensitivity (DEFAULTS)

The ADP approximation depends on the choice of DEFAULTS completion vector. We test 6 alternatives: the actual defaults, 5 other configurations (including the worst known from grid search and the empirically optimal). If the gap $f^* - f_{\text{DP-HPO}}$ is small and consistent, the method is robust to this hyperparameter.

A3. Seed Sensitivity

We vary the random seed from 0 to 24 and report the coefficient of variation ($\text{CV} = \text{std}/\text{mean}$) for each method. DP-HPO's deterministic structure should produce lower CV than stochastic baselines.

A4. Scale Experiment

We extend the search space from 108 to up to 5000 configurations (5–7 dimensions, adding dropout, batch size, weight decay). DP-HPO's evaluation count grows linearly ($O(\Sigma |H_i|)$) while grid search grows exponentially ($O(\Pi |H_i|)$). We plot evaluation reduction ratio and accuracy gap as functions of search space size.

A5. Interaction Strength Measurement

We empirically measure ϵ on each dataset using the estimator from theory.py: $\hat{\epsilon} = \max_{i \neq j} [\max_{\{h_i, h_j\}} |f(h_i, h_j, \text{DEFAULT}_{-ij}) - f(\text{DEFAULT}_i, \text{DEFAULT}_j, \text{DEFAULT}_{-ij})|]$. We report the predicted gap bound $(d-1) \cdot \hat{\epsilon}$ and compare it to the empirically observed gap $f^* - f_{\text{DP-HPO}}$.

A6. Wall-Clock Time Analysis

We report wall-clock time for all 9 methods. DP-HPO and Random $k=10$ use the fewest evaluations; Hyperband and BOHB may use more wall-clock time even at the same trial count due to multi-fidelity overhead. Grid search serves as the absolute upper bound.

IX. DISCUSSION

A. When Does DP-HPO Excel?

DP-HPO is most effective when: (1) training a single model is expensive (large dataset, many epochs), so the 10-vs-108 evaluation difference translates to wall-clock savings; (2) hyperparameter dimensions have low interaction strength ϵ ; (3) the first few dimensions (importance-ordered) dominate performance variation, as is typical for learning rate and architecture depth.

B. Limitations

DP-HPO is designed for small-to-medium discrete search spaces. For continuous or high-dimensional spaces ($d \gg 10$ or $|H_i| \gg 10$), evaluation savings decrease relative to Bayesian methods. The gap bound scales with d and ϵ ; in spaces with high pairwise interactions, DP-HPO may underperform BOHB or Bayesian Optimisation. The DEFAULTS vector must be chosen before running; Ablation A2 quantifies this sensitivity.

C. Extension to Larger Spaces

DP-HPO can be extended by: (1) applying Lemma 2 ordering to focus the first sweeps on the most important dimensions; (2) using a beam rather than a greedy commit at each stage (beam-DP-HPO); (3) adapting the DEFAULTS vector between stages using the committed values so far. These extensions are left for future work.

X. CONCLUSION

We have presented DP-HPO, a framework that formalises neural network HPO as a finite-horizon MDP and solves it via approximate dynamic programming with evaluation caching. Theorem 1 shows DP-HPO is globally optimal under conditional independence and bounded otherwise, with $\text{gap} \leq (d-1) \cdot \epsilon$. Lemma 2 provides a principled importance-first ordering. Proposition 1 proves exactly 10 evaluations suffice for the standard 4-dimensional MLP space, versus 108 for grid search (90.7% reduction). Empirically, DP-HPO remains within 0.5% of grid-search accuracy across all four benchmark datasets, with the accuracy gap not statistically significant on UCI Breast Cancer and statistically significant but practically small on the larger datasets.

We benchmark DP-HPO against eight baselines across four datasets with 25 seeds and full Bonferroni-corrected statistical testing. Across all four datasets, DP-HPO remains within 0.5% of exhaustive grid search accuracy while using only 10 evaluations versus 108—a 90.7% reduction in computational cost. On UCI Breast Cancer the gap is not statistically significant ($p = 0.0036$ at $\alpha_{\text{adj}} = 0.00156$); on the larger datasets (MNIST, Fashion-MNIST, Adult Income) grid search is statistically superior, confirming that DP-HPO is most advantageous when training cost is high and a small accuracy trade-off is acceptable. Code is available at: <https://github.com/kartikjsonawane/dp-hpo>.

REFERENCES

- [1] M. Feurer and F. Hutter, "Hyperparameter Optimization," in *Automated Machine Learning*, Springer, 2019, pp. 3–38.
- [2] J. Bergstra and Y. Bengio, "Random Search for Hyper-Parameter Optimization," *JMLR*, vol. 13, pp. 281–305, 2012.
- [3] B. Shahriari et al., "Taking the Human Out of the Loop: A Review of Bayesian Optimization," *Proc. IEEE*, vol. 104, no. 1, pp. 148–175, 2016.
- [4] L. Li et al., "Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization," *JMLR*, vol. 18, pp. 1–52, 2017.
- [5] S. Falkner, A. Klein, and F. Hutter, "BOHB: Robust and Efficient Hyperparameter Optimization at Scale," *ICML*, 2018, pp. 1437–1446.
- [6] F. Hutter, H. Hoos, and K. Leyton-Brown, "Sequential Model-Based Optimization for General Algorithm Configuration," *LION*, 2011, pp. 507–523.
- [7] J. Bergstra et al., "Algorithms for Hyper-Parameter Optimization," *NeurIPS*, 2011, pp. 2546–2554.
- [8] F. Hutter, H. Hoos, and K. Leyton-Brown, "An Efficient Approach for Assessing Hyperparameter Importance," *ICML*, 2014, pp. 754–762.
- [9] J. N. van Rijn and F. Hutter, "Hyperparameter Importance Across Datasets," *KDD*, 2018, pp. 2367–2376.

- [10] T. Elsken, J. H. Metzen, and F. Hutter, "Neural Architecture Search: A Survey," JMLR, vol. 20, pp. 1–21, 2019.
- [11] W. H. Wolberg, W. N. Street, and O. L. Mangasarian, "Breast Cancer Wisconsin (Diagnostic) Data Set," UCI ML Repository, 1995.
- [12] Y. LeCun, C. Cortes, and C. J. Burges, "The MNIST Database of Handwritten Digits," 1998.
- [13] H. Xiao, K. Rasul, and R. Vollgraf, "Fashion-MNIST: A Novel Image Dataset for Benchmarking Machine Learning Algorithms," arXiv:1708.07747, 2017.
- [14] R. Kohavi, "Scaling Up the Accuracy of Naive-Bayes Classifiers: A Decision-Tree Hybrid," KDD, 1996, pp. 202–207.
- [15] F. Wilcoxon, "Individual Comparisons by Ranking Methods," Biometrics Bulletin, vol. 1, no. 6, pp. 80–83, 1945.
- [16] J. Cohen, Statistical Power Analysis for the Behavioral Sciences, 2nd ed. Lawrence Erlbaum, 1988.
- [17] J. Demšar, "Statistical Comparisons of Classifiers over Multiple Data Sets," JMLR, vol. 7, pp. 1–30, 2006.
- [18] R. Bellman, Dynamic Programming. Princeton University Press, 1957.
- [19] D. P. Bertsekas, Approximate Dynamic Programming, 2nd ed. Athena Scientific, 2012.

APPENDIX: VERSION 1 RESULTS (HISTORICAL REFERENCE)

△ V1 used $n = 5$ seeds. Wilcoxon test cannot achieve $p < 0.05$ at $n = 5$ (min achievable $p = 0.0625$). These results are not statistically valid; shown for reproducibility only.

Method	Mean Acc (%)	Std (%)	Evals	p-value*
Grid Search	98.95	0.35	108	—
Random Search (k=20)	98.77	0.43	20	0.438
Bayesian Opt	98.77	0.43	20	0.438
Optuna (TPE)	98.77	0.43	~16	0.438
DP-HPO	98.60	0.70	10	—

TABLE A-I: UCI Breast Cancer V1 results ($n = 5$ seeds). Not statistically significant.